

njtat7de9

August 1, 2023

```
[7]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
[8]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[9]: df_train=pd.read_csv("/content/drive/MyDrive/mydatasets/C8_loan-train.csv")
df_test=pd.read_csv("/content/drive/MyDrive/mydatasets/C8_loan-test.csv")
```

```
[10]: df_train.dropna(inplace=True)
df_test.dropna(inplace=True)
```

```
[11]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 480 entries, 1 to 613
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Loan_ID               480 non-null   object
 1   Gender                480 non-null   object
 2   Married               480 non-null   object
 3   Dependents            480 non-null   object
 4   Education             480 non-null   object
 5   Self_Employed         480 non-null   object
 6   ApplicantIncome       480 non-null   int64
 7   CoapplicantIncome     480 non-null   float64
 8   LoanAmount            480 non-null   float64
 9   Loan_Amount_Term      480 non-null   float64
10   Credit_History        480 non-null   float64
11   Property_Area         480 non-null   object
```

```
12  Loan_Status      480 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 52.5+ KB
```

```
[12]: df_train.columns
```

```
[12]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
          'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
          'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
         dtype='object')
```

```
[13]: feature_matrix = df_train[['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',  
                                'Loan_Amount_Term', 'Credit_History']]  
target_vector = df_train[['Self_Employed']]
```

```
[14]: fs = StandardScaler().fit_transform(feature_matrix)
      logr = LogisticRegression()
      logr.fit(fs,target_vector)
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)

```

```
[14]: LogisticRegression()
```

```
[15]: observation = df_test[['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',  
                             'Loan_Amount_Term', 'Credit_History']]  
prediction = logr.predict(observation)  
print(prediction)
```

[illegible]

```
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes']
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has
feature names, but LogisticRegression was fitted without feature names
warnings.warn(
```

```
[16]: logr.classes_
```

```
[16]: array(['No', 'Yes'], dtype=object)
```

```
[17]: logr.predict_proba(observation)[0][0]
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has
feature names, but LogisticRegression was fitted without feature names
warnings.warn(
```

```
[17]: 0.0
```